



Fault Detection and Prediction in Models: Optimizing Resource Usage in Cloud Infrastructure

Wei Wu

Amazon

Seattle, WA, USA

wwvela@amazon.com

Abstract

Optimizing resource consumption in cloud infrastructure, the key monitoring, detection and prediction of faults are most relevant task. Optimizing resource usage in the cloud infrastructure is quite challenging. We propose a framework that utilizes Prediction and Monitoring techniques to detect fault and optimize the resource allocation and reduces the downtimes. By employing machine learning algorithms to scrutinise performance data from the past, the framework is able to anticipate fault events before they happen. Besides, this prediction ability takes action at the right moment to improve the operational reliability. A recursive is also built-in, allowing the model to improve performance continuously, based on real-time insights from cloud operations. From the experiments conducted in different cloud environments, considerable enhancements have been obtained in fault detection rates and resource utilization efficiency as compared to those in existing approaches thus highlighting the importance of predictive analytics in building robust open cloud infrastructure.

CCS Concepts

• **Software and its engineering** → **Formal software verification**.

Keywords

Fault Detection, Cloud Infrastructure, Resource Usage

ACM Reference Format:

Wei Wu. 2025. Fault Detection and Prediction in Models: Optimizing Resource Usage in Cloud Infrastructure. In *2025 4th International Conference on Intelligent Systems, Communications and Computer Networks (ISCCN 2025)*, February 21–23, 2025, Nanjing, China. ACM, New York, NY, USA, 8 pages. <https://doi.org/10.1145/3732945.3732964>

1 Introduction

Optimizing resource consumption in cloud infrastructure, the key monitoring, detection and prediction of faults are most relevant task. Advanced methodologies allow for more accurate early defect detection, a key factor in sustaining operational efficiency. Indeed, it has been emphasized that early alerts in resource-intensive environments are key [16], where a multivariate statistical process

control (SPC) method is introduced in [16] and used to extract Fourier transform features to improve early detection of faults in rotating equipment.

In addition, power grid optimization framework characterizes multiple agents with the related failure scenarios, proposing a multi-class prediction model that can predict failures and utilizes these predictions to minimize resource allocation [18]. Furthermore, AI-based methods in arc fault detection are more effective than intelligent systems [29] to optimize the use of resources. Uncertainty-aware models help substantially when used together for fault detection, as they provide a level of certainty in the degree of correctness of fault identification, which is of paramount importance in cloud-based environments, where poor decision making can have severe effects on resource consumption and overall expenses [10]. The amalgamation of statistical methodologies, AI techniques, and uncertainty-aware frameworks presented here exemplifies a roadmap toward more efficient and responsive fault detection mechanisms, essential to cloud management.

But optimizing resource usage in the cloud infrastructure is quite challenging. Giving quantum algorithms, like the quantum approximate optimization algorithm (QAOA), the fault tolerance they need poses a serious problem as error rates can derail such computations. As demonstrated in [13], the Iceberg code is an example of an error detection code that can improve performance through prediction and calibration of expected results by the network. Particularly by identifying and mitigating faults in models, a significant step towards improving efficiency as well as reliability overall. Nevertheless, current techniques might not fully cover the predictive side of fault detection, preventing proper resource optimization in cloud settings. This is a very disciplined path where the development of powerful faults prediction and resource utilization models is essential.

We present a new Framework for Fault Detection and Prediction in Models to address the troublesome drawback of resources being inefficiently utilized in Cloud Infrastructure. This involves the use of proactive monitoring methods that enable the early detection of impending failures in cloud services. Our framework employs machine learning algorithms that analyzes historical performance metrics to alert on potential failures. By being able to pinpoint issues beforehand, they can address the issues addressed before affecting production, greatly decreasing downtime and optimizing resource allocation. In addition, we build a recursive mechanism that iteratively improves model accuracy and efficacy with respect to actionable insights from real-time operation. Through extensive experiments over a wide variety of cloud environments, we show that our framework significantly improve the fault detection rates



This work is licensed under a Creative Commons Attribution International 4.0 License.

ISCCN 2025, Nanjing, China

© 2025 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-1520-4/25/02

<https://doi.org/10.1145/3732945.3732964>

while also increasing the efficiency of resource utilization comparing with traditional way. Results from the research shed light on the role of predictive analytics in empowering better resiliency and resourcefulness in the management of cloud infrastructures.

Our Contributions. Our work has the following key contributions.

- New framework for Fault Detection and Prediction using machine learning to proactively mine cloud infrastructure usage data to detect service failure before it happens.
- By enabling proactive measures before failures happen, our predictions drastically reduce downtime and improve resource management in the cloud.
- We propose a recursive that progressively enhances model accuracy based on real-time operational insights, illustrating substantial improvements in fault detection rates and resource utilization efficiency over state-of-the-art approaches.

2 Related Work

2.1 Anomaly Detection in Cloud System

IBM Cloud presented a new high-dimensional dataset for support of testing methods for anomaly detection with real-world data, essential for the maintenance of large-scale cloud infrastructures [14]. The use of long short-term memory models on a cloud-native data platform for offshore radio access networks to overcome connectivity challenges is among the latest techniques for anomaly detection [1]. Furthermore, a hybrid approach that combines cloud computing and edge computing for edge networks most significantly improves the recall score and F1-score for industrial sensor networks [33]. RCInvestigator: Enables improved analytics of anomaly causes to enhance human-machine collaboration [21]. Other tools include the DeCorus tool, which surpasses previous log fraud detectors in performance [30]. The KAN-AD method improves upon traditional time series anomaly detection by emphasizing global trends that capture the context of anomalies while minimizing noise generated by more localized variations [35]. More recent works have focused on enhancing log-based anomaly detection using a trie structure that accommodates evolving data [20]. We distinguish LPC-AD system as speedy and accurate for multivariate time series anomaly detection in several datasets [26]. Hierarchical approaches are Taught in schools but in a manner that can provide conditional random fields as a much better approach for modelling the dynamics of complex systems, therefore accurately defining their anomalies [24].

2.2 Techniques for Resource Optimization

Resource optimization techniques are becoming increasingly popular in several fields. MicroOpt is another new introduction by Sulaiman et al. that leverages a neural network-based model for smart and effective resource allocation in a 5G network and demonstrates significant improvements in performance [28]. On the other hand, methods as the ones proposed in [17] take advantage of model parallelism to increase the capacity of resource usage still maintaining high accuracy for IoT malware detection. In satellite communication, joint semantic and channel encoding based method reduces the latency and improves transparency effectively [12]. An auction-based adaptive allocation strategy is employed which optimizes

the allocation of communication resources in dense IoT networks with the goal of enabling reliable communication at minimal energy expenditure [31]. Such methods for optimal server resource allocations within data centers have relied on deep reinforcement learning algorithms that take into consideration several relevant parameters while focussing on fault tolerance aspects [7]. As stated in [4], the discovery of resource-efficient large language models leads to tons of methods to improve the performance of these models at the cost of usability. Further, growing models with scalable multilingual feedback can promote cross-multinational adaptation of language models [9]. Large language models focused on the compiler optimization also helps to propel research and development, while also ensuring cost efficiency [8]. Eventually, the gap in performance optimization for serverless applications can be gained by comprehensive resource management and operational efficiency [5] and the low resource NLP model improvement via pruning and distillation emphasizes on the efficiency/accuracy trade-off [23].

2.3 Cloud Infrastructure Predictive Maintenance

An advanced predictive algorithm has been developed to optimize infrastructure partitioning in monolithic software when deployed on microservices infrastructure. This method effectively analyzes dependencies and computational load, enabling a more efficient decomposition of monolithic architectures into microservices. By improving service isolation, fault tolerance, and load balancing, the approach significantly enhances the scalability and maintainability of cloud-based applications. This innovation opens new research directions in microservices architecture, particularly in the areas of distributed cloud networks, dynamic service discovery, and inter-service communication efficiency [25].

Additionally, a new intelligent approach integrating multi-agent Deep Reinforcement Learning (DRL) has been introduced for optimizing time and cost-efficient scaling of serverless applications. By dynamically allocating computing resources based on real-time workload predictions, this method minimizes cold start latency while maintaining function performance balance. The reinforcement learning agents are trained to predict demand fluctuations and proactively adjust scaling policies, ensuring efficient resource utilization without compromising application responsiveness. This approach successfully addresses the challenges of function scaling in serverless environments, paving the way for adaptive, self-optimizing cloud systems [22].

3 Methodology

Our framework for Fault Detection and Prediction addresses this by employing proactive monitoring to detect impending failures. This is achieved by using machine learning algorithms on historical performance data to predict faults in advance, preventing needless downtime and better schedule resources. Real-time operational data allows the model to leverage recursive, improving most of the accuracy over time. We conduct experiments across multiple cloud environments to show that our proposed method can achieve higher fault detection and improve resource utilization over existing methods, clearly indicating the necessity of predictive analytics in the cloud infrastructure.




Figure 1: Fault detection on the maintenance workflow.

3.1 Predictive Fault Monitoring

The Fault Detection and Prediction in Models framework integrates a proactive monitoring approach $Mfdp$. By using the historical performance data D_{hist} , we perform machine learning models MML to predict possible faults F_{pred} . It can be described as:

The historical performance data D_{hist} are collected and preprocessed for analysis. Training the Model: We train a ML model on this data to identify patterns that are indicative of faults. We can summarize the learning objective of the model in as follow:

$$\mathcal{L}(\theta) = \arg \min_{\theta} \sum_{i=1}^n C(f(x_i; \theta), y_i) \quad (1)$$

where $f(x; \theta)$ is the predictive function, y_i are the true labels, and C is a selected cost function.

After the model has been trained, it uses real-time data D_{real} to predict possible faults:

$$F_{pred} = f(D_{real}; \theta) \quad (2)$$

While making predictions F_{pred} , a recursive that compares them with actual outcomes F_{actual} is taken into account. The feedback is used to tweak the model parameters θ according to performance, iteratively improving the system when a loss function:

$$\theta_{new} = \theta_{old} - \eta \nabla \mathcal{L}(F_{pred}, F_{actual}) \quad (3)$$

where η is the learning rate. The framework comprises data-driven continuous learning leveraging operational data to adaptively fine-tune the model, enhancing fault detection rates as well as optimizing resource allocation by reducing downtime. Such predictive fault monitoring is essential to bolster resilience in the management of cloud infrastructure.

3.2 Resource Optimization

We use a machine learning model $\mathcal{M}$ based on previous performance data $\mathcal{D}$ to optimize resource utilization on cloud infrastructure and to avoid potential faults. The prediction model can be mathematically defined as:

$$\mathcal{M}(\mathcal{D}) = P(y|x) \quad (4)$$

where x denotes input features pertaining to the system performance metrics, and y signifies an indication of the subsequent occurrence of a fault. This model trains on data up until October 2023, and using these patterns from $\mathcal{D}$, it gives signals on when and where are the probable faults.

Specifically, to manage resource allocation effectively, we devise an optimization function $\mathcal{R}$ that seeks to minimize resource wastage while ensuring service quality is met. This can be represented as:

$$\min_r \alpha \cdot w_r + \beta \cdot t_{down} \quad (5)$$

We also introduce the feedback function $\mathcal{F}$ to adjust/change the parameters of this model according to its real-time performance by taking performance evaluations in the entire community, which will help fine-tune the predictive capabilities. This update can be expressed mathematically as:

$$\mathcal{F}_{update} = \mathcal{M}_{new}(\mathcal{D}_{real-time}) \quad (6)$$

By leveraging both predictive analytics and continuous recursives to adapt resources quickly, this mechanism improves the resilience and efficiency of resource allocation in cloud environments.

3.3 Machine Learning Deployment

Given a dataset $\mathcal{D}$ of features x_i and their corresponding labels y_i , the learning problem can be formalized as:

$$\hat{y} = f(x_i; \theta) \quad (7)$$

with f being the machine learning model, parameterized by θ . We apply different algorithms to learn the mapping from inputs (the covariates) to outputs (the expected performance). The loss function $\mathcal{L}$ is defined as:

$$\mathcal{L}(\hat{y}, y) = \frac{1}{N} \sum_{i=1}^N \ell(y_i, \hat{y}_i) \quad (8)$$

where ℓ is a particular loss function, and N is the number of examples in the dataset. To obtain our model, we use a training-validation split to ensure the model generalizes well, and we optimize performance by tuning the hyperparameters. Specifically, by utilizing

methods such as cross-validation, you guarantee that your model will teach itself the relevant fault patterns.

Secondly, a continual recursive is established to enhance model accuracy by fine-tuning the model parameters with real-time operational insights. This recursive can be expressed mathematically as:

$$\theta_{new} = \theta_{old} - \eta \nabla \mathcal{L} \quad (9)$$

with η being the learning rate and $\nabla \mathcal{L}$ the gradient of the loss function w.r.t. the model parameters. Thus, supporting both failure localization and evolution of performance numbers in a time-sensitive manner, our framework can reveal potential failures, while enabling it to catch up with evolving performance metrics in the cloud infrastructure management.

4 Experimental Setup

4.1 Datasets

These datasets include NovelCraft (multimodal episodic data for novelty detection learning) [11], LogiQA (logical reasoning in machine reading comprehension) [19], and Pano3D (a benchmark for depth estimation) [2]. The datasets from Acharya et al. may also open avenues for out-of-context and cross-domain adaptations experiments. Lastly, the hierarchical text generation might shed light on strategic dialogue insights in case of human-model interactions in such systems [34].

4.2 Models

Then, we combine them by configuring the Transformer-based models for fault detection and prediction while optimizing for resource usage in cloud infrastructure environments. In particular, we leverage the BERT model (*bert-base-uncased*) to encode operational logs and system metrics into feature vectors. The model, combined with a LightGBM classifier to make predictions, guarantees effective anomaly detection in the datasets and potential resource wasting. We experiments on a dataset whose records include historical usage patterns as well as fault occurrences, allowing us to decide hyper parameters that either improve our models' accuracy or save computational resources. When evaluating the performance of both models with respect to each other, we adopt precision, recall, and F1-score as key metrics to determine their effectiveness when implemented in online environments.

4.3 Implementations

First, we preprocessing the operational logs and system metrics for feature generation through the BERT model which accepts a maximum input length of 512 tokens. In the training phase, we choose a learning rate of $2e-5$ and use the Adam optimizer to achieve better convergence. We set the batch size to 32 which is a compromise between computational need and memory usage while the model is training. For BERT model training, we set the number of epochs to 5 and apply early stopping if no improvement in validation loss is observed over 3 epochs to ensure proper dataset learning.

Evaluating using a 5-fold cross-validation, we are able to provide robust performance metrics on our framework as measured by the metrics of precision, recall, and F1-score as it pertains to

the application of our framework towards our fault detection and resource allocation. We leverage the LightGBM classifier with hyperparameters tuned for our use-case after processing the features. This prevents too many leaves with a max depth equal to 5 to avoid overfitting as well. We still set the value for the learning rate of LightGBM at 0.1.

5 Experiments

5.1 Main Results

Table 1 summarizes the performance metrics of different methods based on BERTs for the cloud infrastructure fault detection and prediction problem.

The BERT + LightGBM model exhibits great strengths in multiple approaches. Out of various methods examined, the highest precision at **0.85** is achieved by the EPST approach reflecting its robustness in predicting positive instances correctly. It also records a decent recall score of **0.80**, so after five training epochs, with learning rate of $2e-5$, it outperforms on F1-score to **0.82**. It reflects the model's balanced ability capturing true positives, while minimizing false negatives.

The EML achieves a precision of **0.82** and recall of **0.76** with an F1-score of **0.79**. Unlike Coefficient Estimation which has precision and recall of **0.80** and **0.78** respectively, leading to an equal performance profile (an F1-score of **0.79**). Both of these use a precision and recall tradeoff while also using five epochs for training at a learning rate of $2e-5$ which suggests good performance.

Precision Bug Identification using Subgraph-Oriented testing achieves an F1-score of **0.78**, along with Precision and Recall of **0.83** and **0.74** respectively, showing that it is predominantly dependent on precision when it comes to finding bugs. This performance shows the method as a useful application of precision-sensitive environments at the cost of a slight recall degradation.

In conclusion, an overview looking at these methods highlights the preforming nature of the BERT + LightGBM pipeline towards our fault detection problem. These different results highlight the opportunities for resource prediction that can improve fault detection rates and resource utilization efficiency in cloud infrastructure using predictive analytics.

5.2 Variations on different architectures

We also explore the impact of each modification in the presented Fault Detection and Prediction in Models as far as the performance of these key metrics. Different setups of the BERT + Modified LightGBM schema, shown in Table 2, are employed in the analysis.

Baseline Without recursive: This configuration employs event based prediction mechanisms without the recursive. It achieves a Precision of 0.83 and a Recall of 0.76 and an F1-score of 0.79 over 5 epochs with a learning rate of $2e-5$, which we'll be using as a baseline for comparison.

- **Baseline Without recursive:** This configuration employs event based prediction mechanisms without the recursive. It achieves a Precision of 0.83 and a Recall of 0.76 and an F1-score of 0.79 over 5 epochs with a learning rate of $2e-5$, which we'll be using as a baseline for comparison.
- **Thresholded Feature Set:** Here, we build a conditional model for smaller feature set coefficient estimation. This results in

Table 1: Evaluation metrics for several methods with respect to F1-score, Epochs, and Learning Rate. The BERT framework is employed for fault detection in cloud infrastructure based on each method.

Model	F1-score	Epochs	Learning Rate	Method	Precision	Recall
<i>BERT + LightGBM</i>						
EPST [3]	0.82	5	2e-5	Event Pred.	0.85	0.80
C-SDE [6]	0.79	5	2e-5	Coeff. Est.	0.80	0.78
EML [27]	0.79	5	2e-5	Ens. Meth.	0.82	0.76
SOT [32]	0.78	5	2e-5	Prec. Bug ID	0.83	0.74
VA-P [15]	0.80	5	2e-5	Pred. Maint. Strat.	0.86	0.75

a set of metrics with Precision = 0.78, Recall = 0.72 and an F1-score of 0.75, both with 5 epochs and a learning rate of 2e-5.

- *Model with more epochs:* Using some ensemble methods well as after 20 train epochs, this model produced the best results yet of 0.84 precision, 0.77 recall, and 0.80 f1 score. The best learning rate is 1e-5.
- *Non about Predictive Feedback:* In this experiment, we design a precision bug diagnosis method without the predictive feedback component. The model shows Precision = 0.80, Recall = 0.70 and F1-score = 0.75 after 5 epochs at 2e-5.
- *Venn-Abers without Adjustment Mechanism:* The metrics under predictive maintenance strategies with no adjustment mechanism show Precision=0.82, Recall=0.71, F1-score=0.76, all measured across 5 epochs and 2e-5 learning rate.

Table 2 shows the results of the ablation studies which highlight the importance of each part of our model in detecting faults. In summary, the recursive timing and reducing epoch definitely show better performance metrics, indicating that the use of successive model iterations and adaptive methods can significantly improve detection rates and the efficient use of resources in a cloud environment. Every configuration provides insight as to how certain structural components influence the overall predictive power of the framework thus proving that well-thought-out modifications can induce drastic effect in the fault identification and prediction performance in cloud landscapes.

5.3 Proactive Monitoring Techniques

Cloud infrastructure faults detection cannot be solely based upon post-incident analysis and only can greatly be optimized with proactive monitoring techniques. The comparative performance of different techniques as explored in Table 3.

Predictive analytics shows higher effectiveness in fault finding. Predictive analytics is the most effective, with a detection rate of 0.85, indicating its strengths in identifying potential failures at an early stage compared to other methods. The false positive rate is lowest at 0.05 in this technique suggestive of better reliability in detecting the issue rather than producing unprecedented false alerts. Moreover, predictive analytics leads to considerably less latency at 95 ms compared to competitors, allowing faster fault response.

But traditional approaches such as threshold-based monitoring, anomaly detection, and log analysis have had mixed success. The




Figure 2: Analysis of historical performance data demonstrating the effect of various methods on downtime reduction, accuracy, resource utilization.

monitoring based on thresholds offers a detection rate of 0.75 while its false-positive rate is 0.10 and its latency is 120 ms, and the anomaly detection techniques are a bit better with a 0.80 detection rate and a false-positive rate of 0.08 (still not capable of surpassing the efficiency done by predictive analytics). According to log analysis, this model shows more balanced performance with detection rate of 0.78, however latency is significantly higher at 150 ms.

Effective cloud management relies on resource utilization

While being the most effective at detection, predictive analytics also returns a strong resource utilization percentage of 72% versus 70% in threshold-based monitoring, 65% in anomaly detection, and 68% in log analysis. But these metrics from the paper demonstrate the advantage of combining predictive analytics with cloud infrastructure management, as it can help improve both fault detection and resource management efficiency.

5.4 Data Analysis on Previous Information

By analyzing historical performances data, we can determine how effective certain approaches are in the improvement of cloud infrastructure management. The performance metrics of the several approaches against accuracy, resource utilization, and downtime mitigation are illustrated in Figure 2.

Venn-Abers outperforms on all metrics evaluated. With an accuracy rate of 0.90, Venn-Abers also achieved the highest

Table 2: Results of ablation studies assessing the influence of different component exclusions and modifications on metrics for carelessness detection in cloud infrastructure. The results from both methods illustrate the change on F1-score, Epochs, and Learning Rate while dealing with the BERT frame.

Model	F1-score	Epochs	Learning Rate	Ablation Method	Precision	Recall
Baseline w/o FB Loop	0.79	5	2e-5	Event Pred.	0.83	0.76
Baseline w/ Red. Feat.	0.75	5	2e-5	Coeff. Est.	0.78	0.72
Model w/ More Epochs	0.80	20	1e-5	Ens. Meth.	0.84	0.77
Cond. w/o Pred. FB	0.75	5	2e-5	Prec. Bug ID	0.80	0.70
Venn-Abers w/o Adj. Mech.	0.76	5	2e-5	Pred. Maint. Strat.	0.82	0.71

Table 3: Comparison of proactive monitoring techniques for fault detection in cloud infrastructure, including detection rate, resource utilization, false positives, and latency.

Tech.	Det. Rate	Res. Util.	FP	Lat. (ms)
Thresh. Mon.	0.75	70%	0.10	120
Anom. Det. Alg.	0.80	65%	0.08	100
Log Anal. Meth.	0.78	68%	0.12	150
Pred. Analyt.	0.85	72%	0.05	95

resource utilization (0.93) and reduced downtime with 32%. It indicates its strength in detecting and predicting faults enabling optimal resource management.

EPST and Ensemble ML are a close second. In contrast, EPST stands at 0.88 for accuracy and 0.92 for the resource utilization metric with 30% lesser downtime and therefore a good alternative. Ensemble ML provides an accuracy of 0.86 along with a resource utilization of 0.91 with 28% downtime reduction proving its efficacy too.

CT for SDE and Testing on Subgraph are slightly lowered performance. Controlled SDE accuracy=0.84; Controlled SDE resource utilization=0.90; Downtime=25% reduced. Results by Subgraph Testing achieve accuracy of 0.85, resource usage at 0.88, and a reduced downtime of 26%, with these approaches still enjoying to have some gaps with respect to best methods proposed.

Finally, the comparison of the performance results highlights that several innovative frameworks are suitable to predict faults and optimize resource usage in cloud infrastructures, validating the support of the machine learning techniques on proactive monitoring and management.

5.5 Machine Learning Algorithm Implementation

Different metrics are defined for different algorithms, and the test results show that Nearest Neighbors was effective even with standard metrics when used on these small, hyper-parameter sets across cloud infrastructure fault detection and prediction with various machine learning approaches.

A baseline for fault detection metrics comes from traditional algorithms. Logistic Regression (Table 4) performs with an accuracy of 0.75, an F1-score of 0.73, and a ROC AUC of 0.76. It

Table 4: Comparative analysis of machine learning algorithms for cloud infrastructure fault detection Well-known metrics.

Model	Algorithm	Accuracy	F1-score	AUC
<i>Advanced Methods</i>				
XGBoost	XGBoost	0.85	0.83	0.84
Neural Network	Deep Learning	0.88	0.87	0.90
SVM	Support Vector Machine	0.84	0.82	0.85
<i>Traditional Methods</i>				
Logistic Regression	Logistic Reg.	0.75	0.73	0.76
Decision Tree	Decision Tree	0.78	0.76	0.79
Random Forest	Random Forest	0.80	0.79	0.81

looks like the decision tree with maximal depth of 3 and a minimum of samples per leaf (base case) teaches the machine fairly well, with random forest slightly outperforming it (increased max depth, more trees, more target classes): it achieves 0.80 accuracy, 0.79 f1-score and 0.81 ROC AUC. Despite establishing a solid foundation, these models may fail to demonstrate their abilities under complex fault scenarios, which are common in cloud services.

Hierarchical algorithms improve the ability to detect faults. These advanced methods show significant improvements on prediction metrics. XGBoost achieves an accuracy of 0.85, an F1-score of 0.83 and a ROC AUC of 0.84. Also, we see an improvement when using the Neural Network model, as it reaches an accuracy of 0.88, precision of 0.87 and 0.90 of ROC AUC, showing that it is the one with the greatest ability to process complex faults. The SVM also evaluates well with 0.84 accuracy and 0.82 F1-score indicating that such models retain a balanced performance to some extent compared to classical approaches.

The results underline that with the increase in the factors affecting cloud system usage, machine learning is becoming more effective in managing cloud infrastructure fault detection. By leveraging the power of sophisticated machine learning algorithms, we not only report heightened detection rates but also improved prediction processes leading to better resource allocation and operational efficiencies in the cloud ecosystem.

5.6 Recursive Mechanism

The recursive component is an important factor to refine the fault detection and prediction frame. Introducing feedback from Real




Figure 3: Performance of various techniques in optimizing fault detection through the recursive mechanism, highlighting improvements in accuracy, reliability, and system efficiency.

Time Operational Data, to greatly improve model accuracy and efficiency. The analysis shown in Figure 3 is built on the evaluation of individual techniques in terms of performance optimization.

The integration of this recursive mechanism effectively enhances the detection of faults by 15%, decreases the operating time by 20% and raises its exactness by 10%, thus projects the ability of the framework to learn and adjust its forecasts in real-time dependent on performance or on various criteria of performance, and takes a proactive approach to resource consumption.

Real-time Data Analysis enhances response. Real-time data analysis is where it all starts, providing an 18% average increase in detection deterministic. What is more, it also achieves a 25% time savings for fault finding and a 12% gain in accuracy. This approach enables the framework to stay reactive to real-time shifts in the state of the cloud environment, thus facilitating prompt detection and resolution of faults.

Among the three proposed methods, the predicted adjustment method has the most dramatic effect, achieving 20% improvement over the fault detection effectiveness at the overall level. Moreover, it reduces operations by 30% and improves accuracy by 15%. Overall, this method emphasizes how powerful using historical data would enhance predicting future outcomes, allowing services to run efficiently and avoid disturbances.

6 Limitations

Limitations of Proposed Model Fault Detection and Prediction in Models First of all, since it is trained on historical performance, if there is no history, atypical situation or new fault type, it can face troubles. This may hinder its ability to predict in real-time as new patterns arise. Moreover, feedback, i.e. using the model as a recursive, might only improve the accuracy of the model if operational data is relevant and of sufficient quality, given the lack of training after October 2023. In case the data is inconsistent or have noise, it can hurt the model's performance. Future lines of work could explore including more adaptive techniques or hybrid methods that decrease dependency on historical data, thus

providing more robustness under out of ordinary conditions and increasing efficiency on the fly.

7 Conclusions

This method also incorporates active monitoring, which enables the timely discovery of impending service outages by utilizing machine learning algorithms to mine historical performance information for potential malfunctions. This predictive capability also facilitates timely intervention, reducing downtime and optimizing resource allocation. Moreover, the architecture incorporates a recursive that regularly improves model precision by learning from operational data in real time. Our experimental results across diverse cloud computing environments demonstrate that utilizing this fault prediction framework maximizes fault detection rates and effectively improves resource utilization compared to traditional techniques, highlighting the importance of predictive analytics in advancing fault management within cloud infrastructures.

References

- [1] A. Ahmad, Peizheng Li, Robert J. Piechocki, and Rui Inacio. 2024. Anomaly Detection in Offshore Open Radio Access Network Using Long Short-Term Memory Models on a Novel Artificial Intelligence-Driven Cloud-Native Data Platform. *ArXiv abs/2409.02849* (2024).
- [2] G. Albanis, N. Zioulis, Petros Drakoulis, V. Gkitsas, V. Sterzentsenko, Federico Álvarez, D. Zarpalas, and P. Daras. 2021. Pano3D: A Holistic Benchmark and a Solid Baseline for 360° Depth Estimation. *2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)* (2021), 3722–3732.
- [3] Evie Andrew, Travis Monk, and A. Schaik. 2023. An Event based Prediction Suffix Tree. *ArXiv abs/2310.14944* (2023).
- [4] Guangji Bai, Zheng Chai, Chen Ling, Shiyu Wang, Jiaying Lu, Nan Zhang, Tingwei Shi, Ziyang Yu, Mengdan Zhu, Yifei Zhang, Carl Yang, Yue Cheng, and Liang Zhao. 2024. Beyond Efficiency: A Systematic Survey of Resource-Efficient Large Language Models. *ArXiv abs/2401.00625* (2024).
- [5] Mohamed Lemine El Bechir, Cheikh Sad Bouh, and Abobakr Shuwail. 2024. Comprehensive Review of Performance Optimization Strategies for Serverless Applications on AWS Lambda. *ArXiv abs/2407.10397* (2024).
- [6] Luc Brogat-Motte, Riccardo Bonalli, and Alessandro Rudi. 2024. Learning Controlled Stochastic Differential Equations. *ArXiv abs/2411.01982* (2024).
- [7] Chang-Lin Chen, Hanhan Zhou, Jiayu Chen, Mohammad Pedramfar, Vaneet Aggarwal, Tian Lan, Zheqing Zhu, Chi Zhou, Tim Gasser, Pol Mauri Ruiz, Vijay Menon, Neeraj Kumar, and Hongbo Dong. 2023. Two-tiered Online Optimization of Region-wide Datacenter Resource Allocation via Deep Reinforcement Learning. *ArXiv abs/2306.17054* (2023).
- [8] Chris Cummins, Volker Seeker, Dejan Grubisic, Baptiste Rozière, Jonas Gehring, Gabriele Synnaeve, and Hugh Leather. 2024. Meta Large Language Model Compiler: Foundation Models of Compiler Optimization. *ArXiv abs/2407.02524* (2024).
- [9] John Dang, Arash Ahmadian, Kelly Marchisio, Julia Kreutzer, A. Ustun, and Sara Hooker. 2024. RLHF Can Speak Many Languages: Unlocking Multilingual Preference Optimization for LLMs. (2024), 13134–13156.
- [10] L. Das, B. Gjorgiev, and G. Sansavini. 2023. Uncertainty-aware deep learning for digital twin-driven monitoring: Application to fault detection in power lines. *ArXiv abs/2303.10954* (2023).
- [11] Patrick Feeney, Sarah Schneider, Panagiotis Lymperopoulos, Liping Liu, Matthias Scheutz, Michael C. Hughes Dept. of Computer Science, T. University, Center for Vision, Automation, Control, and Austrian Institute of Technology. 2022. NovelCraft: A Dataset for Novelty Detection and Discovery in Open Worlds. *Trans. Mach. Learn. Res.* 2023 (2022).
- [12] Sheikh Salman Hassan, Loc X. Nguyen, Y. Tun, Zhu Han, and Choong Seon Hong. 2024. Semantic Enabled 6G LEO Satellite Communication for Earth Observation: A Resource-Constrained Network Optimization. *ArXiv abs/2408.03959* (2024).
- [13] Zichang He, David Amaro, Ruslan Shaydulin, and Marco Pistoia. 2024. Performance of Quantum Approximate Optimization with Quantum Error Detection. *ArXiv abs/2409.12104* (2024).
- [14] Mohammad Saiful Islam, M. Rakha, William Pourmajidi, Janakan Sivaloganathan, John Steinbacher, and Andriy V. Miransky. 2024. Anomaly Detection in Large-Scale Cloud Systems: An Industry Case and Dataset. *ArXiv abs/2411.09047* (2024).
- [15] U. Johansson, Tuwe Lofstrom, and C. Sonstrod. 2023. Well-Calibrated Probabilistic Predictive Maintenance using Venn-Abers. *ArXiv abs/2306.06642* (2023).

- [16] V. Jorry, Zina-Sabrina Duma, T. Sihvonen, S. Reinikainen, and L. Roininen. 2024. Statistical Batch-Based Bearing Fault Detection. *ArXiv abs/2407.17236* (2024).
- [17] Sreenitha Kasarapu, Sanket Shukla, and Sai Manoj Pudukotai Dinakarrao. 2024. Enhancing IoT Malware Detection through Adaptive Model Parallelism and Resource Optimization. *ArXiv abs/2404.08808* (2024).
- [18] Malte Lehna, Mohamed Hassouna, Dmitry Degtyar, Sven Tomforde, and Christoph Scholz. 2024. Fault Detection for agents on power grid topology optimization: A Comprehensive analysis. *ArXiv abs/2406.16426* (2024).
- [19] Jian Liu, Leyang Cui, Hanmeng Liu, Dandan Huang, Yile Wang, and Yue Zhang. 2020. LogiQA: A Challenge Dataset for Machine Reading Comprehension with Logical Reasoning. *ArXiv abs/2007.08124* (2020).
- [20] Jinyang Liu, Junjie Huang, Yintong Huo, Zhihan Jiang, Jiazhen Gu, Zhuangbin Chen, Cong Feng, Minzhi Yan, and Michael R. Lyu. 2023. Scalable and Adaptive Log-based Anomaly Detection with Expert in the Loop. *ArXiv abs/2306.05032* (2023).
- [21] Shuhan Liu, Yunfan Zhou, Lu Ying, Yuan Tian, Jue Zhang, Shandan Zhou, Weiwei Cui, Qingwei Lin, T. Moscibroda, Haidong Zhang, Di Weng, and Yingcai Wu. 2024. RCInvestigator: Towards Better Investigation of Anomaly Root Causes in Cloud Computing Systems. *ArXiv abs/2405.15571* (2024).
- [22] Anupama Mampage, S. Karunasekera, and R. Buyya. 2023. A Deep Reinforcement Learning based Algorithm for Time and Cost Optimized Scaling of Serverless Applications. *ArXiv abs/2308.11209* (2023).
- [23] Aishwarya Mirashi, Purva Lingayat, Srushti Sonavane, Tejas Padhiyar, Raviraj Joshi, and Geetanjali Kale. 2024. On Importance of Pruning and Distillation for Efficient Low Resource NLP. *ArXiv abs/2409.14162* (2024).
- [24] Srishti Mishra, Tvarita Jain, and D. Sitaram. 2022. A Hierarchical Approach to Conditional Random Fields for System Anomaly Detection. *ArXiv abs/2210.15030* (2022).
- [25] Kalyani Pendyala and R. Buyya. 2024. An Infrastructure Cost Optimised Algorithm for Partitioning of Microservices. *ArXiv abs/2408.06570* (2024).
- [26] Zhi Qi, Hong Xie, Ye Li, Jian Tan, Feifei Li, and John C.S. Lui. 2022. LPC-AD: Fast and Accurate Multivariate Time Series Anomaly Detection via Latent Predictive Coding. *ArXiv abs/2205.08362* (2022).
- [27] Sheikh Md. Mushfiqur Rahman and Nasir U. Eisty. 2024. Introducing Ensemble Machine Learning Algorithms for Automatic Test Case Generation using Learning Based Testing. *ArXiv abs/2409.04651* (2024).
- [28] Muhammad Sulaiman, Mahdieh Ahmadi, Bo Sun, Niloy Saha, M. A. Salahuddin, R. Boutaba, and Aladdin Saleh. 2024. MicroOpt: Model-driven Slice Resource Optimization in 5G and Beyond Networks. *ArXiv abs/2407.18342* (2024).
- [29] Kriti Thakur, Divyanshi Dwivedi, K. V. Sam, Moses Babu, Pradeep Alivelu Manga Parimi, Kumar Yemula, Pratyush Chakraborty, M. Pal, Ms. Kriti Thakur, Mrs. Divyanshi Dwivedi, Mr. K. Victor Sam, and Dr. Pradeep Kumar Yemula. 2023. Advancements in Arc Fault Detection for Electrical Distribution Systems: A Comprehensive Review from Artificial Intelligence Perspective. *ArXiv abs/2311.16804* (2023).
- [30] Bruno Wassermann, David Ohana, Ronen Schaffer, Robert J. Shahla, E. K. Kolodner, Eran Raichstein, and Michal Malka. 2022. DeCorus: Hierarchical Multivariate Anomaly Detection at Cloud-Scale. *ArXiv abs/2202.06892* (2022).
- [31] Nirmal D. Wickramasinghe, John Dooley, Dirk Pesch, and Indrakshi Dey. 2024. Auction-based Adaptive Resource Allocation Optimization in Dense IoT Networks. *ArXiv abs/2409.17843* (2024).
- [32] Xiaoyuan Xie, Yan Song, Songqiang Chen, and Jinfu Chen. 2024. Subgraph-Oriented Testing for Deep Learning Libraries. *ArXiv abs/2412.06430* (2024).
- [33] Tao Yang, Jinming Wang, Weijie Hao, Qiang Yang, and Wen-Hung Wang. 2022. Hybrid Cloud-Edge Collaborative Data Anomaly Detection in Industrial Sensor Networks. *ArXiv abs/2204.09942* (2022).
- [34] Denis Yarats and M. Lewis. 2017. Hierarchical Text Generation and Planning for Strategic Dialogue. *ArXiv abs/1712.05846* (2017).
- [35] Quan Zhou, Changhua Pei, Fei Sun, Jing Han, Zhengwei Gao, Dan Pei, Haiming Zhang, Gaogang Xie, and Jianhui Li. 2024. KAN-AD: Time Series Anomaly Detection with Kolmogorov-Arnold Networks. *ArXiv abs/2411.00278* (2024).